

# Can You Predict Weight Lifting Mistakes from Accelerometer Data?

by Scott Semel

January 28, 2015

## Executive Summary

Six young health participants were asked to perform one set of 10 repetitions of the Unilateral Dumbbell Biceps Curl in five different fashions: exactly according to the specification (Class A), throwing the elbows to the front (Class B), lifting the dumbbell only halfway (Class C), lowering the dumbbell only halfway (Class D) and throwing the hips to the front (Class E). Class A corresponds to the specified execution of the exercise, while the other 4 classes correspond to common mistakes. We trained and tested several classifiers to find out that accurate prediction of the Class type was possible depending on the method with between 70 and 97% accuracy.

Read more: [http://groupware.les.inf.puc-rio.br/har#weight\\_lifting\\_exercises#ixzz3yZxNTUIP](http://groupware.les.inf.puc-rio.br/har#weight_lifting_exercises#ixzz3yZxNTUIP)

## Data Processing

```
load("workspace2016-01-26")
# The Weightlifting data set as 19622 records with 160
columns.
library(AppliedPredictiveModeling)
library(caret)

## Loading required package: lattice

## Loading required package: ggplot2

dat = read.csv("pml-training.csv")
test = read.csv("pml-testing.csv")
# Several columns in the test set has only NA or mostly NA so
we removed those.
test2 = test[ , ! apply( test , 2 , function(x) any(is.na(x)) )
]
# And we realized it was only necessary to train on those
columns in the test set. There were more columns in the
training set that would never get used anyway.
keepvar = names(test2)
keepvar = keepvar[-60]
keepvar = c(keepvar, "classe")
# It was important to also remove the first 8 columns from the
```

training and test set because there were counters and timestamps that we did not want to affect the outcome. Leaving those in confused the classifiers which thought they were doing better than they were.

```
keepvar = keepvar[8:60]
dat2 = dat[, keepvar]
# We trained on 75% and tested on the other 25%. It did not
# appear to matter where the seed started and still got very
# close solutions.
inTrain = createDataPartition(y=dat2$classe, p = .75,
list=FALSE)
training = dat2[ inTrain,]
testing = dat2[-inTrain,]
# The linear discriminant analysis did not work as well only
# producing 70% accuray
# modFit1 = train(classe~ .,data=training, method="lda")
pred = predict(modFit1,testing)
```

```
## Loading required package: MASS
```

```
#confusionMatrix(pred, testing$classe)
predict(modFit1,test)
```

```
## [1] B A B C C E D D A A D A B A E A A B B B
## Levels: A B C D E
```

*# We tried it with and without cross-validation and it only increased accuracy slightly. Also performing the centering and standardizing between 0 and 1 did not make a difference in the prediction so we didn't continue to do that.*

```
tc = trainControl("repeatedcv",
                  number=10,
                  repeats=10,
                  classProbs=TRUE,
                  savePred=T)
# modFit2 = train(classe~ .,data=training,
method="lda",trControl=tc,
#               preProc=c("center", "scale"))
pred = predict(modFit2,testing)
#table(pred, testing$classe)
#confusionMatrix(pred, testing$classe)
predict(modFit2,test)
```

```
## [1] B A B C C C D D A A D A B A E A A B B B
## Levels: A B C D E
```

*# The boosted trees method worked with high accuracy around 97%  
# modFit3 = train(classe~ .,data=training, method="gbm")  
# It was helpful to see the importance of each variable. At first we tried very small subsets of important variables. Then we realized we could test all of them. It is interesting that*

*it has different lists of importance each time so we couldnt just use this function to tell us which variables to use.*

```
gbmImp = varImp(modFit3, scale = FALSE)

## Loading required package: gbm

## Loading required package: survival

##
## Attaching package: 'survival'

## The following object is masked from 'package:caret':
##
##      cluster

## Loading required package: splines

## Loading required package: parallel

## Loaded gbm 2.1.1

## Loading required package: plyr

pred = predict(modFit3,testing)
confusionMatrix(pred, testing$classe)

## Confusion Matrix and Statistics
##
##              Reference
## Prediction      A      B      C      D      E
##      A 1378    25      0      0      3
##      B   14   908    26      0      6
##      C    1    15   812    32      6
##      D    1     1    13   769    18
##      E    1     0     4     3   868
##
## Overall Statistics
##
##              Accuracy : 0.9655
##              95% CI : (0.96, 0.9705)
##      No Information Rate : 0.2845
##      P-Value [Acc > NIR] : < 2.2e-16
##
##              Kappa : 0.9564
##      McNemar's Test P-Value : 0.0001159
##
## Statistics by Class:
##
##              Class: A Class: B Class: C Class: D
Class: E
## Sensitivity              0.9878    0.9568    0.9497    0.9565
0.9634
```

```
## Specificity          0.9920    0.9884    0.9867    0.9920
0.9980
## Pos Pred Value      0.9801    0.9518    0.9376    0.9589
0.9909
## Neg Pred Value      0.9951    0.9896    0.9894    0.9915
0.9918
## Prevalence          0.2845    0.1935    0.1743    0.1639
0.1837
## Detection Rate      0.2810    0.1852    0.1656    0.1568
0.1770
## Detection Prevalence 0.2867    0.1945    0.1766    0.1635
0.1786
## Balanced Accuracy   0.9899    0.9726    0.9682    0.9742
0.9807
```

```
predict(modFit3,test)
```

```
## [1] B A B A A E D B A A B C B A E E A B B B
## Levels: A B C D E
```

```
# The rpart method did only produced around 65% accuracy for
some reason.
```

```
# modFit4 = train(classe~ .,data=training, method="rpart",)
pred = predict(modFit4,testing)
```

```
## Loading required package: rpart
```

```
#confusionMatrix(pred, testing$classe)
predict(modFit4,test)
```

```
## [1] C D D C C C D D A A C D C A D D C D C D
## Levels: A B C D E
```

```
# The random forest method worked the best with possibly more
than 97% accuracy.
```

```
# modFit5 = train(classe~ .,data=training, method="rf")
pred = predict(modFit5,testing)
```

```
## Loading required package: randomForest
```

```
## randomForest 4.6-12
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##
```

```
## Attaching package: 'randomForest'
```

```
## The following object is masked from 'package:ggplot2':
```

```
##
```

```
## margin
```

```
confusionMatrix(pred, testing$classe)
```

```

## Confusion Matrix and Statistics
##
##           Reference
## Prediction   A     B     C     D     E
##           A 1395     0     0     0     0
##           B     0  949     0     0     0
##           C     0     0  855     3     0
##           D     0     0     0  801     2
##           E     0     0     0     0  899
##
## Overall Statistics
##
##           Accuracy : 0.999
##           95% CI : (0.9976, 0.9997)
##           No Information Rate : 0.2845
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9987
##           Mcnemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D
Class: E
## Sensitivity           1.0000    1.0000    1.0000    0.9963
0.9978
## Specificity           1.0000    1.0000    0.9993    0.9995
1.0000
## Pos Pred Value        1.0000    1.0000    0.9965    0.9975
1.0000
## Neg Pred Value        1.0000    1.0000    1.0000    0.9993
0.9995
## Prevalence            0.2845    0.1935    0.1743    0.1639
0.1837
## Detection Rate        0.2845    0.1935    0.1743    0.1633
0.1833
## Detection Prevalence  0.2845    0.1935    0.1750    0.1637
0.1833
## Balanced Accuracy      1.0000    1.0000    0.9996    0.9979
0.9989

predict(modFit5,test)

## [1] B A B A A E D B A A B C B A E E A B B B
## Levels: A B C D E

# The next step would be to combine all the models and either
average the answers or vote with them. But in this case random
forests scored 100% on the 20 question validation set so it was
not necessary.

```

## Results

A relationship was found between the class and the predictor variables.

We successfully reduced the number of useful variables to these:

```
gbmImp
```

```
## gbm variable importance
##
##   only 20 most important variables shown (out of 52)
##
##               Overall
## roll_belt      2907.6
## pitch_forearm  1529.4
## yaw_belt       1256.5
## magnet_dumbbell_z  940.0
## magnet_dumbbell_y  802.7
## roll_forearm    733.5
## magnet_belt_z    567.2
## gyros_belt_z     453.4
## roll_dumbbell    387.4
## pitch_belt      370.7
## accel_forearm_x  367.4
## gyros_dumbbell_y  321.6
## accel_dumbbell_y  286.1
## magnet_belt_y    228.5
## magnet_forearm_z  211.2
## accel_dumbbell_x  206.7
## yaw_arm         188.4
## accel_belt_z     167.6
## magnet_dumbbell_x 148.7
## magnet_arm_z     144.3
```

Even from this small data set the work appears promising.

Velloso, E.; Bulling, A.; Gellersen, H.; Ugulino, W.; Fuks, H. Qualitative Activity Recognition of Weight Lifting Exercises. Proceedings of 4th International Conference in Cooperation with SIGCHI (Augmented Human '13) . Stuttgart, Germany: ACM SIGCHI, 2013.

Read more: [http://groupware.les.inf.puc-rio.br/har#weight\\_lifting\\_exercises#ixzz3yZx0m9BL](http://groupware.les.inf.puc-rio.br/har#weight_lifting_exercises#ixzz3yZx0m9BL)